

# Joining Tables

## Contents

<b>Baseball Data</b>	<b>1</b>
<b>Teams data</b>	<b>2</b>
<b>Salaries data</b>	<b>4</b>
<b>Join Teams and Salaries</b>	<b>5</b>

Data is generally stored in multiple smaller tables. These smaller tables share a common variable or field, thus are related (which is why these tables are also known as relations). All the relations together comprise a relational database.

Storing data across relations offers a number of benefits such as reduced redundancy, fewer errors from updates, and greater efficiency in storage and retrieval. The data required for analysis at hand is retrieved by linking relations through joins. The languages most widely used for this is SQL, however, this note will illustrate the joins using R.

## Baseball Data

To illustrate joins, we will use baseball data compiled in library(Lahman) based on the Lahman Database

```
library(Lahman)
```

Baseball data is spread across a number of smaller tables.

```
library(printr)
data(package = "Lahman")
```

Table 1: Data sets in Lahman

Item	Title
AllstarFull	AllstarFull table
Appearances	Appearances table
AwardsManagers	AwardsManagers table
AwardsPlayers	AwardsPlayers table
AwardsShareManagers	AwardsShareManagers table
AwardsSharePlayers	AwardsSharePlayers table
Batting	Batting table
BattingPost	BattingPost table
CollegePlaying	CollegePlaying table
Fielding	Fielding table
FieldingOF	FieldingOF table
FieldingOFsplit	FieldingOFsplit table
FieldingPost	FieldingPost data
HallOfFame	Hall of Fame Voting Data
HomeGames	HomeGames table

Item	Title
LahmanData	Lahman Datasets
Managers	Managers table
ManagersHalf	ManagersHalf table
Master	Master table
Parks	Parks table
People	People table
Pitching	Pitching table
PitchingPost	PitchingPost table
Salaries	Salaries table
Schools	Schools table
SeriesPost	SeriesPost table
Teams	Teams table
TeamsFranchises	TeamFranchises table
TeamsHalf	TeamsHalf table
battingLabels	Variable Labels
fieldingLabels	Variable Labels
pitchingLabels	Variable Labels

## Teams data

In a well-designed relational database, each table has a variable, called a primary key, that uniquely identifies each row. Let us see if this is the case for the `Teams` table.

```
detach('package:printr', unload = TRUE)
head(Teams)
```

```
##   yearID lgID teamID franchID divID Rank  G Ghome  W  L DivWin WCWin LgWin
## 1  1871   NA   BS1      BNA  <NA>   3 31   NA 20 10  <NA> <NA>   N
## 2  1871   NA   CH1      CNA  <NA>   2 28   NA 19  9  <NA> <NA>   N
## 3  1871   NA   CL1      CFC  <NA>   8 29   NA 10 19  <NA> <NA>   N
## 4  1871   NA   FW1      KEK  <NA>   7 19   NA  7 12  <NA> <NA>   N
## 5  1871   NA   NY2      NNA  <NA>   5 33   NA 16 17  <NA> <NA>   N
## 6  1871   NA   PH1      PNA  <NA>   1 28   NA 21  7  <NA> <NA>   Y
##   WSWin  R  AB  H X2B X3B HR BB SO SB CS HBP SF  RA  ER  ERA CG SHO SV
## 1 <NA> 401 1372 426  70  37  3 60 19 73 16  NA NA 303 109 3.55 22  1  3
## 2 <NA> 302 1196 323  52  21 10 60 22 69 21  NA NA 241  77 2.76 25  0  1
## 3 <NA> 249 1186 328  35  40  7 26 25 18  8  NA NA 341 116 4.11 23  0  0
## 4 <NA> 137  746 178  19  8  2 33  9 16  4  NA NA 243  97 5.17 19  1  0
## 5 <NA> 302 1404 403  43  21  1 33 15 46 15  NA NA 313 121 3.72 32  1  0
## 6 <NA> 376 1281 410  66  27  9 46 23 56 12  NA NA 266 137 4.95 27  0  0
##   IPouts  HA HRA BBA SOA  E DP  FP      name
## 1   828 367  2  42  23 243 24 0.834  Boston Red Stockings
## 2   753 308  6  28  22 229 16 0.829  Chicago White Stockings
## 3   762 346 13  53  34 234 15 0.818  Cleveland Forest Citys
## 4   507 261  5  21  17 163  8 0.803  Fort Wayne Kekiongas
## 5   879 373  7  42  22 235 14 0.840  New York Mutuals
## 6   747 329  3  53  16 194 13 0.845  Philadelphia Athletics
##                                     park attendance BPF PPF teamIDBR teamIDlahman45
## 1                South End Grounds I          NA 103  98      BOS      BS1
## 2          Union Base-Ball Grounds          NA 104 102      CHI      CH1
## 3 National Association Grounds          NA  96 100      CLE      CL1
## 4                Hamilton Field          NA 101 107      KEK      FW1
```

```
## 5      Union Grounds (Brooklyn)      NA 90 88      NYU      NY2
## 6      Jefferson Street Grounds      NA 102 98      ATH      PH1
##      teamIDretro
## 1      BS1
## 2      CH1
## 3      CL1
## 4      FW1
## 5      NY2
## 6      PH1
```

Let us check to see if teamID uniquely identifies each row.

```
library(dplyr)
Teams %>%
  group_by(teamID, name)%>%
  summarize(n = n())%>%
  filter(n>1)%>%
  arrange(desc(n))

## # A tibble: 116 x 3
## # Groups:   teamID [90]
##   teamID name      n
##   <fct> <chr>    <int>
## 1 PIT    Pittsburgh Pirates    130
## 2 PHI    Philadelphia Phillies  129
## 3 CIN    Cincinnati Reds      125
## 4 SLN    St. Louis Cardinals   121
## 5 CHA    Chicago White Sox     120
## 6 DET    Detroit Tigers        120
## 7 CHN    Chicago Cubs          118
## 8 BOS    Boston Red Sox        113
## 9 NYA    New York Yankees      108
## 10 CLE   Cleveland Indians     106
## # ... with 106 more rows

head(Teams)

##   yearID lgID teamID franchID divID Rank  G Ghome  W  L DivWin WCWin LgWin
## 1  1871  NA   BS1      BNA  <NA>    3 31   NA 20 10  <NA>  <NA>    N
## 2  1871  NA   CH1      CNA  <NA>    2 28   NA 19  9  <NA>  <NA>    N
## 3  1871  NA   CL1      CFC  <NA>    8 29   NA 10 19  <NA>  <NA>    N
## 4  1871  NA   FW1      KEK  <NA>    7 19   NA  7 12  <NA>  <NA>    N
## 5  1871  NA   NY2      NNA  <NA>    5 33   NA 16 17  <NA>  <NA>    N
## 6  1871  NA   PH1      PNA  <NA>    1 28   NA 21  7  <NA>  <NA>    Y
##   WSWin  R  AB  H X2B X3B HR BB SO SB CS HBP SF  RA  ER  ERA  CG SHO SV
## 1  <NA> 401 1372 426 70 37 3 60 19 73 16  NA NA 303 109 3.55 22  1  3
## 2  <NA> 302 1196 323 52 21 10 60 22 69 21  NA NA 241 77 2.76 25  0  1
## 3  <NA> 249 1186 328 35 40 7 26 25 18  8  NA NA 341 116 4.11 23  0  0
## 4  <NA> 137 746 178 19  8 2 33 9 16  4  NA NA 243 97 5.17 19  1  0
## 5  <NA> 302 1404 403 43 21 1 33 15 46 15  NA NA 313 121 3.72 32  1  0
## 6  <NA> 376 1281 410 66 27 9 46 23 56 12  NA NA 266 137 4.95 27  0  0
##   IPouts HA HRA BBA SOA  E DP  FP      name
## 1    828 367  2  42  23 243 24 0.834  Boston Red Stockings
## 2    753 308  6  28  22 229 16 0.829  Chicago White Stockings
## 3    762 346 13  53  34 234 15 0.818  Cleveland Forest Citys
## 4    507 261  5  21  17 163  8 0.803  Fort Wayne Kekiongas
```

```
## 5      879 373    7 42 22 235 14 0.840      New York Mutuals
## 6      747 329    3 53 16 194 13 0.845 Philadelphia Athletics
##           park attendance BPF PPF teamIDBR teamIDlahman45
## 1      South End Grounds I      NA 103 98      BOS      BS1
## 2      Union Base-Ball Grounds      NA 104 102      CHI      CH1
## 3 National Association Grounds      NA 96 100      CLE      CL1
## 4      Hamilton Field      NA 101 107      KEK      FW1
## 5      Union Grounds (Brooklyn)      NA 90 88      NYU      NY2
## 6      Jefferson Street Grounds      NA 102 98      ATH      PH1
## teamIDretro
## 1      BS1
## 2      CH1
## 3      CL1
## 4      FW1
## 5      NY2
## 6      PH1
```

It is clear that `teamID` is repeated multiple times, so it does not uniquely identify each row. From reviewing the first few rows of `Teams`, you have probably realized that the table lists teams for each Year. So, let us see if each row is uniquely defined by `teamID` and `yearID`.

```
Teams %>%
  group_by(teamID, yearID)%>%
  summarize(n = n())%>%
  filter(n>1)

## # A tibble: 0 x 3
## # Groups:   teamID [0]
## # ... with 3 variables: teamID <fct>, yearID <int>, n <int>
```

## Salaries data

From the above, it is clear that `teamID` and `yearID` uniquely identify each row in `Teams`. While this table is a rich source of information on team performance, it doesn't contain any information on player Salary. Salary data is contained in a table called `Salaries`. However, salary data is at a more granular level, providing compensation for each player.

```
head(Salaries)

##   yearID teamID lgID  playerID salary
## 1   1985    ATL   NL barkele01 870000
## 2   1985    ATL   NL bedrost01 550000
## 3   1985    ATL   NL benedbr01 545000
## 4   1985    ATL   NL campri01 633333
## 5   1985    ATL   NL ceronri01 625000
## 6   1985    ATL   NL chambch01 800000
```

We will roll up the data to the team level by computing average player salary.

```
team_salary =
  Salaries %>%
  group_by(yearID, teamID)%>%
  summarize(avgSalary = mean(salary, na.rm=T))
team_salary

## # A tibble: 918 x 3
## # Groups:   yearID [32]
```

```
##   yearID teamID avgSalary
##   <int> <fct>    <dbl>
## 1  1985 ATL      673045.
## 2  1985 BAL      525487.
## 3  1985 BOS      435902.
## 4  1985 CAL      515282.
## 5  1985 CHA      468866.
## 6  1985 CHN      577405.
## 7  1985 CIN      379996.
## 8  1985 CLE      327583.
## 9  1985 DET      517407.
## 10 1985 HOU      499653.
## # ... with 908 more rows
```

## Join Teams and Salaries

You will note that the `team_salary` data by `teamID` and `yearID` just like the `Teams` table. So, we can join the `Teams` data with `Salaries` using `teamID` and `yearID`. There are four kinds of joins: inner join, left outer, right outer and full outer join. An inner join will keep rows for which the keys match across the two tables. A left outer join will keep all rows in the table in the left and matching rows from the right. A right outer join will keep all rows in the table in the right and matching rows from the left. Finally, a full outer join will keep rows that match across the two tables and the ones that don't match.

We want to only keep data for which there is a match across the `Teams` and `Salaries` table, so we implement an inner join using the `dplyr` function `inner_join`. An alternative is to use `merge()` from Base R.

```
Teams %>%
  inner_join(team_salary, by = c('yearID' = 'yearID', 'teamID' = 'teamID'))%>%
  head()
```

##	yearID	lgID	teamID	franchID	divID	Rank	G	Ghome	W	L	DivWin	WCWin	LgWin
## 1	1985	NL	ATL	ATL	W	5	162	81	66	96	N	<NA>	N
## 2	1985	AL	BAL	BAL	E	4	161	81	83	78	N	<NA>	N
## 3	1985	AL	BOS	BOS	E	5	163	81	81	81	N	<NA>	N
## 4	1985	AL	CAL	ANA	W	2	162	79	90	72	N	<NA>	N
## 5	1985	AL	CHA	CHW	W	3	163	81	85	77	N	<NA>	N
## 6	1985	NL	CHN	CHC	E	4	162	81	77	84	N	<NA>	N

##	WSWin	R	AB	H	X2B	X3B	HR	BB	SO	SB	CS	HBP	SF	RA	ER	ERA	CG	SHO	SV
## 1	N	632	5526	1359	213	28	126	553	849	72	52	22	41	781	679	4.19	9	9	29
## 2	N	818	5517	1451	234	22	214	604	908	69	43	19	40	764	694	4.38	32	6	33
## 3	N	800	5720	1615	292	31	162	562	816	66	27	30	57	720	659	4.06	35	8	29
## 4	N	732	5442	1364	215	31	153	648	902	106	51	39	35	703	633	3.91	22	8	41
## 5	N	736	5470	1386	247	37	146	471	843	108	56	43	45	720	656	4.07	20	8	39
## 6	N	686	5492	1397	239	28	150	562	937	182	49	18	39	729	666	4.16	20	8	42

##	IPouts	HA	HRA	BBA	SOA	E	DP	FP	name
## 1	4372	1512	134	642	776	159	197	0.976	Atlanta Braves
## 2	4282	1480	160	568	793	129	168	0.979	Baltimore Orioles
## 3	4384	1487	130	540	913	145	161	0.977	Boston Red Sox
## 4	4372	1453	171	514	767	112	202	0.982	California Angels
## 5	4355	1411	161	569	1023	111	152	0.982	Chicago White Sox
## 6	4327	1492	156	519	820	134	150	0.979	Chicago Cubs

##	park	attendance	BPF	PPF	teamIDBR	teamIDlahman45
## 1	Atlanta-Fulton County Stadium	1350137	105	106	ATL	ATL
## 2	Memorial Stadium	2132387	97	97	BAL	BAL
## 3	Fenway Park II	1786633	104	104	BOS	BOS

```
## 4           Anaheim Stadium    2567427 100 100      CAL      CAL
## 5           Comiskey Park      1669888 104 104      CHW      CHA
## 6           Wrigley Field      2161534 110 110      CHC      CHN
##   teamIDretro avgSalary
## 1      ATL  673045.5
## 2      BAL  525486.9
## 3      BOS  435902.4
## 4      CAL  515281.9
## 5      CHA  468865.6
## 6      CHN  577405.3
```

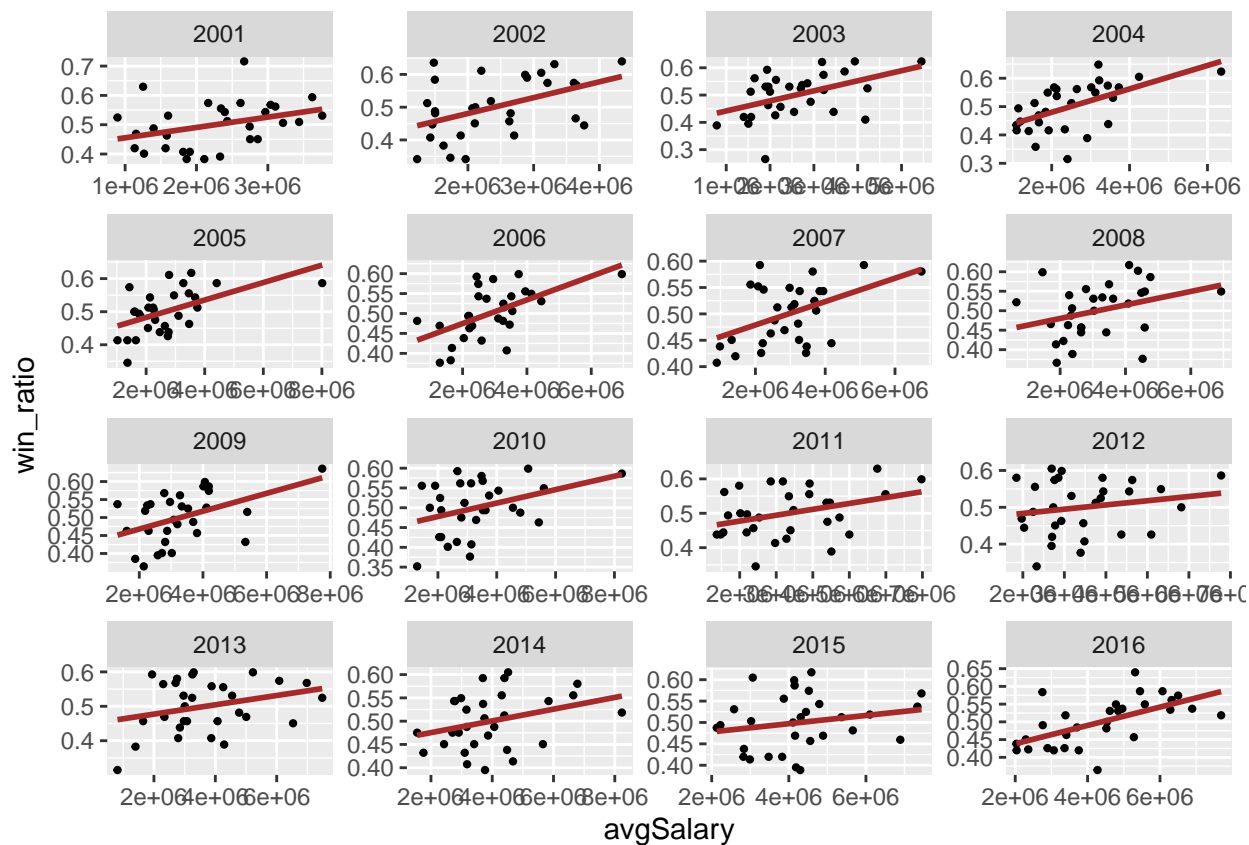
The joined table contains data on both team performance and average team salary. Now, one can filter the data to see if average salary paid is related to good performance or if it draws more fans. We measure performance as proportion of wins.

```
Teams %>%
  inner_join(team_salary, by = c('yearID' = 'yearID', 'teamID' = 'teamID'))%>%
  mutate(win_ratio = W/(W+L))%>%
  select(yearID, teamID, name, win_ratio, avgSalary, attendance)%>%
  head()

##   yearID teamID           name win_ratio avgSalary attendance
## 1  1985    ATL   Atlanta Braves 0.4074074  673045.5    1350137
## 2  1985    BAL Baltimore Orioles 0.5155280  525486.9    2132387
## 3  1985    BOS   Boston Red Sox 0.5000000  435902.4    1786633
## 4  1985    CAL California Angels 0.5555556  515281.9    2567427
## 5  1985    CHA Chicago White Sox 0.5246914  468865.6    1669888
## 6  1985    CHN   Chicago Cubs  0.4782609  577405.3    2161534
```

Let us visualize the relationship between avgSalary and win\_ratio for the the last 16 years. One of the consequences of an inner join is that it will only contain data for years that are present in both tables. Salary data is only available until 2016, so even though Teams data goes past 2016, it is not present in the joined table.

```
library(ggplot2)
Teams %>%
  inner_join(team_salary, by = c('yearID' = 'yearID', 'teamID' = 'teamID'))%>%
  mutate(win_ratio = W/(W+L))%>%
  select(yearID, teamID, name, win_ratio, avgSalary, attendance)%>%
  filter(yearID%in%2001:2016)%>%
  ggplot(aes(x=avgSalary,y=win_ratio))+
  geom_point(size=0.8)+
  geom_smooth(method='lm', se=F,color='brown')+
  facet_wrap(~yearID,scales = 'free')
```



By joining two different tables, we were able to examine the relationship between avgSalary paid to each player and the team win-ratio. What do you think? (Caveat: If you suspect relationship of salary on performance is delayed, then you can examine the relationship between win\_ratio and a lagged avgSalary)