

missing data

Contents

Examine Missing Data	1
Solutions	4
Delete	5
Encode Missing Data	6
Impute	8
mice	8
caret	10

Missing data are quite common in real datasets. This may be driven by a number of reasons including, data was not available, was not gathered, or coded incorrectly.

A number of predictive models require complete data. And, even among the predictive models that can function with less than complete data, missing data can only exist in predictors not the response variable. Before we proceed to find a solution to missing data let us examine the reasons for it. The reasons may be broadly categorized as 1. Random: The cause of missing data may be completely random. Survey respondents may neglect to answer certain questions by chance. Person coding data may commit a data entry error. 2. Data deficiency: This may be a missing component of a predictor. Consider a survey containing the alternatives male and female for a question on gender. Those not conforming to a particular gender may not respond.

3. Specific Causes: There are situations where missing data can be clearly attributed to a cause. Survey respondents may refuse to respond to questions on politically sensitive issues such as abortion and right to bear arms. In clinical studies where patients are measured periodically over time, a patient may drop off due to an adverse side effect.

Examine Missing Data

The pattern of missing data is likely to shed light on the reasons behind it. For small and medium datasets, it is possible to visualize missing data in a chart such as a heatmap. When the dataset has a large number of variables or observations, the data must be suitably condensed before visualizing.

To illustrate, let us create some missing data in the `mtcars` dataset

```
mtcars_missing = mtcars
for(i in 1:35){
  set.seed(i)
  x = sample(1:nrow(mtcars), 1)
  y = sample(1:ncol(mtcars), 1)
  mtcars_missing[x, y] = NA
}
mtcars_missing
```

##	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
## Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	NA
## Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	NA	NA
## Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1


## Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
## Hornet Sportabout	18.7	NA	360.0	175	3.15	3.440	17.02	0	0	NA	2
## Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	NA	3	NA
## Duster 360	14.3	8	NA	245	3.21	3.570	15.84	NA	0	3	4
## Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
## Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	NA	4	2
## Merc 280	NA	NA	NA	123	3.92	NA	18.30	NA	0	4	4
## Merc 280C	17.8	6	167.6	123	3.92	3.440	18.90	1	NA	4	4
## Merc 450SE	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3
## Merc 450SL	NA	8	275.8	180	3.07	3.730	17.60	0	0	3	3
## Merc 450SLC	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3
## Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4
## Lincoln Continental	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4
## Chrysler Imperial	NA	8	440.0	230	3.23	5.345	17.42	0	0	3	4
## Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	NA	1	4	1
## Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
## Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
## Toyota Corona	21.5	4	120.1	97	3.70	NA	20.01	1	0	NA	1
## Dodge Challenger	15.5	8	318.0	150	2.76	3.520	16.87	0	0	3	2
## AMC Javelin	15.2	8	304.0	150	3.15	3.435	17.30	0	0	3	2
## Camaro Z28	13.3	8	NA	245	3.73	3.840	15.41	0	0	3	NA
## Pontiac Firebird	19.2	8	400.0	NA	3.08	3.845	17.05	0	0	3	2
## Fiat X1-9	27.3	NA	79.0	66	4.08	1.935	18.90	1	1	4	1
## Porsche 914-2	26.0	4	120.3	91	NA	2.140	16.70	0	1	5	2
## Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2
## Ford Pantera L	NA	8	351.0	264	4.22	3.170	14.50	NA	1	NA	4
## Ferrari Dino	19.7	6	145.0	175	3.62	2.770	15.50	0	1	5	6
## Maserati Bora	NA	8	301.0	335	3.54	NA	14.60	0	1	5	8
## Volvo 142E	21.4	4	121.0	NA	4.11	2.780	18.60	NA	1	4	2

When the dataset has a reasonable number of observations and variables, a simple method to visualize missing information is a heatmap. With library(ggplot2), `geom_tile` can be used to generate a simple heatmap.

We first classify all observations based on whether they are missing (0) or not (1)

```
mtcars_missing_bi = mtcars_missing
mtcars_missing_bi[!is.na(mtcars_missing)] = 1
mtcars_missing_bi[is.na(mtcars_missing)] = 0

library(tidyr); library(dplyr); library(ggplot2)
mtcars_missing_bi %>%
  as_tibble()%>%
  mutate(obs = rownames(mtcars_missing_bi))%>%
  select(obs, everything())%>%
  pivot_longer(cols = 2:12, names_to = 'variables', values_to = 'values' )%>%
  ggplot(aes(variables, obs, fill= factor(values))) +
  geom_tile()+
  guides(fill=F)
```



Identifying patterns in a moderate sized heatmap can be challenging. To aid interpretation, `library(ComplexHeatMap)` groups variables and observations and represents the groups using a dendrogram. The author has a nice book on the bells and whistles of `ComplexHeatMap`.

`ComplexHeatMap` can be installed from Bioconductor


```
# install.packages('BiocManager')
# BiocManager::install('ComplexHeatmap')
```

Alternatively, the most recent version can be obtained from github

```
#library(devtools)
#install_github("jokergoo/ComplexHeatmap")
```

The following heatmap groups the tiles to make it easier to spot patterns in missing data.

```
library(ComplexHeatmap)
library(circlize)
Heatmap(mtcars_missing_bi,
  rect_gp = gpar(col = "white", lwd = 1),
  name = 'Seeing What is Missing',
  column_title = 'Predictors',
  row_title = 'Observations',
  col = circlize::colorRamp2(c(0,1),c('indianred','aquamarine4')),
  show_heatmap_legend = F,
  row_names_gp = gpar(fontsize = 5))
```



For larger datasets, the data needs to be reduced (using a technique such as principal components analysis) before constructing a heatmap. Another valuable tool for examining the nature and degree of missing data are numerical summaries. For the `mtcars_missing`, for each variable, here are the (a) number of missing values

```
apply(mtcars_missing,
      MARGIN = 2,
      FUN = function(x) sum(is.na(x)))

## mpg cyl disp hp drat wt qsec vs am gear carb
## 5 3 3 2 1 3 0 5 3 4 4
```

(b) percent of missing values

```
apply(mtcars_missing,
      MARGIN = 2,
      FUN = function(x) 100*sum(is.na(x))/(sum(is.na(x))+ sum(!is.na(x))))

## mpg cyl disp hp drat wt qsec vs am gear carb
## 15.625 9.375 9.375 6.250 3.125 9.375 0.000 15.625 9.375 12.500 12.500
```

Solutions

Before we examine solutions, it is worth noting that while many popular predictive models such as linear regression, penalized regression models, support vector machines and neural networks cannot tolerate any amount of missing values, there are a few predictive models such as classification and regression trees and other tree-based models that can handle incomplete data.

For the vast majority of predictive models and also a number unsupervised learning methods, missing values need to be addressed. There are three broad categories of solutions, (a) delete (or ignore) missing values, (b) encode missing data, (c) impute.

Delete

By far, the simplest solution to missing data is remove them either by deleting or ignoring. The most common approach is to conduct listwise deletion wherein the entire row is deleted if even a single element is missing. In `mtcars_missing`, the following rows are missing at least one value.

```
apply(mtcars_missing, MARGIN = 1, function(x) any(is.na(x)))
```

##	Mazda RX4	Mazda RX4 Wag	Datsun 710	Hornet 4 Drive
##	TRUE	TRUE	FALSE	FALSE
##	Hornet Sportabout	Valiant	Duster 360	Merc 240D
##	TRUE	TRUE	TRUE	FALSE
##	Merc 230	Merc 280	Merc 280C	Merc 450SE
##	TRUE	TRUE	TRUE	FALSE
##	Merc 450SL	Merc 450SLC	Cadillac Fleetwood	Lincoln Continental
##	TRUE	FALSE	FALSE	FALSE
##	Chrysler Imperial	Fiat 128	Honda Civic	Toyota Corolla
##	TRUE	TRUE	FALSE	FALSE
##	Toyota Corona	Dodge Challenger	AMC Javelin	Camaro Z28
##	TRUE	FALSE	FALSE	TRUE
##	Pontiac Firebird	Fiat X1-9	Porsche 914-2	Lotus Europa
##	TRUE	TRUE	TRUE	FALSE
##	Ford Pantera L	Ferrari Dino	Maserati Bora	Volvo 142E
##	TRUE	FALSE	TRUE	TRUE

Removing them leaves only 13 of the 32 rows in `mtcars_missing`!

```
mtcars_missing[!apply(mtcars_missing, MARGIN = 1, function(x) any(is.na(x))),]
```

##	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
## Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
## Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
## Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
## Merc 450SE	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3
## Merc 450SLC	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3
## Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4
## Lincoln Continental	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4
## Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
## Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
## Dodge Challenger	15.5	8	318.0	150	2.76	3.520	16.87	0	0	3	2
## AMC Javelin	15.2	8	304.0	150	3.15	3.435	17.30	0	0	3	2
## Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2
## Ferrari Dino	19.7	6	145.0	175	3.62	2.770	15.50	0	1	5	6

Many predictive modeling functions (e.g., `lm`) will automatically conduct listwise deletion. Note in the following output, “19 observations deleted due to missingness”

```
summary(lm(mpg~.,mtcars_missing))
```

```
##
## Call:
## lm(formula = mpg ~ ., data = mtcars_missing)
##
## Residuals:
```

```
##           Datsun 710           Hornet 4 Drive           Merc 240D           Merc 450SE
##           -1.108e+00           -7.352e-01           -6.661e-16           9.495e-01
##           Merc 450SLC Cadillac Fleetwood Lincoln Continental           Honda Civic
##           2.421e-01           2.706e-01           2.980e-02           4.134e-01
##           Toyota Corolla Dodge Challenger           AMC Javelin           Lotus Europa
##           6.943e-01           1.810e+00           -2.566e+00           7.352e-01
##           Ferrari Dino
##           -7.352e-01
##
```

```
## Coefficients:
```

```
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept) -280.98109 166.37953 -1.689 0.233
## cyl         19.61883 13.19206 1.487 0.275
## disp        -0.04405 0.04758 -0.926 0.452
## hp          -0.17579 0.16670 -1.055 0.402
## drat         7.01319 5.20404 1.348 0.310
## wt           8.77079 11.02143 0.796 0.510
## qsec         3.41823 2.04149 1.674 0.236
## vs          26.01966 19.08962 1.363 0.306
## am          10.87792 10.71197 1.015 0.417
## gear        25.82345 15.41268 1.675 0.236
## carb        -3.63447 3.04106 -1.195 0.355
##
```

```
## Residual standard error: 2.683 on 2 degrees of freedom
```

```
## (19 observations deleted due to missingness)
```

```
## Multiple R-squared: 0.9794, Adjusted R-squared: 0.8765
```

```
## F-statistic: 9.518 on 10 and 2 DF, p-value: 0.09876
```

An important consideration in using deletion is loss of sample. With listwise deletion, even a small number of missing elements can result in a large number of rows getting deleted. Consider the above example where the original data was missing only 11.34375% values but listwise deletion eliminated 59.375% of the rows

```
sum(is.na(mtcars_missing))/nrow(mtcars)*ncol(mtcars) # Percent missing values
```

```
## [1] 11.34375
```

```
100* nrow(mtcars_missing[apply(mtcars_missing, MARGIN = 1, function(x) any(is.na(x))),])/nrow(mtcars_mi
```

```
## [1] 59.375
```

A second consideration is the cost of acquiring more observations.

A third and perhaps most important consideration is the likelihood of introducing bias into the model. Consider a clinical study where half the respondents receive the current standard care of treatment while the other half receive the new treatment. The new treatment may induce adverse effects causing patients to drop out of the study. Such missing data is clearly not random and deleting them is bound to bias results.

Encode Missing Data

Missing data may represent information. In a survey that only contains two alternatives for gender, those not conforming to either gender may skip the question. In this case, non-response may contain information valuable gender orientation. As another example, survey respondents may refuse to respond to questions on politically sensitive issues such as abortion and right to bear arms. Here again, non-response may contain useful information.

In a categorical variable, missing data may be encoded as another level. Consider the following example.

```
set.seed(61710)
```

```
gender = sample(c('Male','Female',NA),10,T)
```

```

gender

## [1] NA      "Female" "Male"   NA      "Male"   "Male"   "Female" "Male"
## [9] "Female" "Female"

gender_encoded = gender
gender_encoded[is.na(gender)] = 'Did Not Respond'
gender_encoded

## [1] "Did Not Respond" "Female"      "Male"      "Did Not Respond"
## [5] "Male"           "Male"       "Female"     "Male"
## [9] "Female"         "Female"

```

Missing data in a continuous variable may be represented using a dummy variable.

```

mtcars_missing_encoded = mtcars_missing
mtcars_missing_encoded %>%
  mutate(cyl_missing = ifelse(is.na(mtcars_missing_encoded$cyl),0,1))%>%
  select(mpg, cyl, cyl_missing, everything())

##           mpg  cyl  cyl_missing  disp  hp  drat    wt  qsec  vs  am  gear
## Mazda RX4      21.0    6           1 160.0 110 3.90 2.620 16.46 0  1    4
## Mazda RX4 Wag  21.0    6           1 160.0 110 3.90 2.875 17.02 0  1   NA
## Datsun 710      22.8    4           1 108.0  93 3.85 2.320 18.61 1  1    4
## Hornet 4 Drive  21.4    6           1 258.0 110 3.08 3.215 19.44 1  0    3
## Hornet Sportabout 18.7   NA           0 360.0 175 3.15 3.440 17.02 0  0   NA
## Valiant         18.1    6           1 225.0 105 2.76 3.460 20.22 1 NA    3
## Duster 360      14.3    8           1    NA 245 3.21 3.570 15.84 NA  0    3
## Merc 240D        24.4    4           1 146.7  62 3.69 3.190 20.00 1  0    4
## Merc 230         22.8    4           1 140.8  95 3.92 3.150 22.90 1 NA    4
## Merc 280         NA   NA           0    NA 123 3.92    NA 18.30 NA  0    4
## Merc 280C        17.8    6           1 167.6 123 3.92 3.440 18.90 1 NA    4
## Merc 450SE       16.4    8           1 275.8 180 3.07 4.070 17.40 0  0    3
## Merc 450SL       NA    8           1 275.8 180 3.07 3.730 17.60 0  0    3
## Merc 450SLC      15.2    8           1 275.8 180 3.07 3.780 18.00 0  0    3
## Cadillac Fleetwood 10.4    8           1 472.0 205 2.93 5.250 17.98 0  0    3
## Lincoln Continental 10.4    8           1 460.0 215 3.00 5.424 17.82 0  0    3
## Chrysler Imperial  NA    8           1 440.0 230 3.23 5.345 17.42 0  0    3
## Fiat 128         32.4    4           1  78.7  66 4.08 2.200 19.47 NA  1    4
## Honda Civic      30.4    4           1  75.7  52 4.93 1.615 18.52 1  1    4
## Toyota Corolla   33.9    4           1  71.1  65 4.22 1.835 19.90 1  1    4
## Toyota Corona    21.5    4           1 120.1  97 3.70    NA 20.01 1  0   NA
## Dodge Challenger  15.5    8           1 318.0 150 2.76 3.520 16.87 0  0    3
## AMC Javelin      15.2    8           1 304.0 150 3.15 3.435 17.30 0  0    3
## Camaro Z28       13.3    8           1    NA 245 3.73 3.840 15.41 0  0    3
## Pontiac Firebird  19.2    8           1 400.0  NA 3.08 3.845 17.05 0  0    3
## Fiat X1-9        27.3   NA           0  79.0  66 4.08 1.935 18.90 1  1    4
## Porsche 914-2    26.0    4           1 120.3  91  NA 2.140 16.70 0  1    5
## Lotus Europa     30.4    4           1  95.1 113 3.77 1.513 16.90 1  1    5
## Ford Pantera L   NA    8           1 351.0 264 4.22 3.170 14.50 NA  1   NA
## Ferrari Dino     19.7    6           1 145.0 175 3.62 2.770 15.50 0  1    5
## Maserati Bora    NA    8           1 301.0 335 3.54    NA 14.60 0  1    5
## Volvo 142E       21.4    4           1 121.0  NA 4.11 2.780 18.60 NA  1    4
##           carb
## Mazda RX4      NA
## Mazda RX4 Wag  NA
## Datsun 710      1

```

## Hornet 4 Drive	1
## Hornet Sportabout	2
## Valiant	NA
## Duster 360	4
## Merc 240D	2
## Merc 230	2
## Merc 280	4
## Merc 280C	4
## Merc 450SE	3
## Merc 450SL	3
## Merc 450SLC	3
## Cadillac Fleetwood	4
## Lincoln Continental	4
## Chrysler Imperial	4
## Fiat 128	1
## Honda Civic	2
## Toyota Corolla	1
## Toyota Corona	1
## Dodge Challenger	2
## AMC Javelin	2
## Camaro Z28	NA
## Pontiac Firebird	2
## Fiat X1-9	1
## Porsche 914-2	2
## Lotus Europa	2
## Ford Pantera L	4
## Ferrari Dino	6
## Maserati Bora	8
## Volvo 142E	2

Impute

In imputation, information and relationships among non-missing predictors is used to estimate missing values. Approaches to imputation differ based on whether they are designed for inferences or prediction. Our focus is going to be on the latter, so the quality of an imputation procedure will be judged by its ability to accurately predict missing values.

It is important to remember that imputed values are merely estimates of the true value for missing data. Thus, one must limit how much missing data is imputed. While there isn't a universal rule on how much data can be imputed, no more than 20% for a variable may be a good thumb rule to use.

Imputation must occur prior to other steps in data preparation.

There are a number of methods available for imputing missing data (e.g., predictive mean matching, k-nearest neighbors, trees and tree-based methods) and a number of R package that implement these (e.g., `mice`, `caret`)

`mice`

`mice()` implements a number of imputation methods. This illustration uses the default, which for numeric data is predictive mean matching. For more information, see author's website. In version 3.12.0, `mice` library implemented a new `matchindex C` function that makes predictive mean matching 50 to 600 times faster, however this affects reproducibility of the algorithm. Read more about the update (here)[<https://cran.r-project.org/web/packages/mice/news/news.html>]. (If you wish to reproduce the behavior of the previous version of `mice`, include `use.matcher=T` in the `mice` function). Finally, `tidyr` has an identically named `complete` function. To prevent conflicts, it is best to include the package reference, `mice::complete(...)`


```

library(mice)
mtcars_mice = mice::complete(mice(mtcars_missing, seed = 617))

##
## iter imp variable
## 1 1 mpg cyl disp hp drat wt vs am gear carb
## 1 2 mpg cyl disp hp drat wt vs am gear carb
## 1 3 mpg cyl disp hp drat wt vs am gear carb
## 1 4 mpg cyl disp hp drat wt vs am gear carb
## 1 5 mpg cyl disp hp drat wt vs am gear carb
## 2 1 mpg cyl disp hp drat wt vs am gear carb
## 2 2 mpg cyl disp hp drat wt vs am gear carb
## 2 3 mpg cyl disp hp drat wt vs am gear carb
## 2 4 mpg cyl disp hp drat wt vs am gear carb
## 2 5 mpg cyl disp hp drat wt vs am gear carb
## 3 1 mpg cyl disp hp drat wt vs am gear carb
## 3 2 mpg cyl disp hp drat wt vs am gear carb
## 3 3 mpg cyl disp hp drat wt vs am gear carb
## 3 4 mpg cyl disp hp drat wt vs am gear carb
## 3 5 mpg cyl disp hp drat wt vs am gear carb
## 4 1 mpg cyl disp hp drat wt vs am gear carb
## 4 2 mpg cyl disp hp drat wt vs am gear carb
## 4 3 mpg cyl disp hp drat wt vs am gear carb
## 4 4 mpg cyl disp hp drat wt vs am gear carb
## 4 5 mpg cyl disp hp drat wt vs am gear carb
## 5 1 mpg cyl disp hp drat wt vs am gear carb
## 5 2 mpg cyl disp hp drat wt vs am gear carb
## 5 3 mpg cyl disp hp drat wt vs am gear carb
## 5 4 mpg cyl disp hp drat wt vs am gear carb
## 5 5 mpg cyl disp hp drat wt vs am gear carb

head(mtcars_mice)

##          mpg cyl disp  hp drat   wt  qsec vs am gear carb
## Mazda RX4      21.0   6  160 110 3.90 2.620 16.46 0  1   4    3
## Mazda RX4 Wag  21.0   6  160 110 3.90 2.875 17.02 0  1   4    4
## Datsun 710      22.8   4  108  93 3.85 2.320 18.61 1  1   4    1
## Hornet 4 Drive  21.4   6  258 110 3.08 3.215 19.44 1  0   3    1
## Hornet Sportabout 18.7   8  360 175 3.15 3.440 17.02 0  0   3    2
## Valiant         18.1   6  225 105 2.76 3.460 20.22 1  0   3    1

library(mice)
mtcars_mice_rf = mice::complete(mice(mtcars_missing, method = 'rf', seed = 617))

##
## iter imp variable
## 1 1 mpg cyl disp hp drat wt vs am gear carb
## 1 2 mpg cyl disp hp drat wt vs am gear carb
## 1 3 mpg cyl disp hp drat wt vs am gear carb
## 1 4 mpg cyl disp hp drat wt vs am gear carb
## 1 5 mpg cyl disp hp drat wt vs am gear carb
## 2 1 mpg cyl disp hp drat wt vs am gear carb
## 2 2 mpg cyl disp hp drat wt vs am gear carb
## 2 3 mpg cyl disp hp drat wt vs am gear carb
## 2 4 mpg cyl disp hp drat wt vs am gear carb
## 2 5 mpg cyl disp hp drat wt vs am gear carb

```

```
## 3 1 mpg cyl disp hp drat wt vs am gear carb
## 3 2 mpg cyl disp hp drat wt vs am gear carb
## 3 3 mpg cyl disp hp drat wt vs am gear carb
## 3 4 mpg cyl disp hp drat wt vs am gear carb
## 3 5 mpg cyl disp hp drat wt vs am gear carb
## 4 1 mpg cyl disp hp drat wt vs am gear carb
## 4 2 mpg cyl disp hp drat wt vs am gear carb
## 4 3 mpg cyl disp hp drat wt vs am gear carb
## 4 4 mpg cyl disp hp drat wt vs am gear carb
## 4 5 mpg cyl disp hp drat wt vs am gear carb
## 5 1 mpg cyl disp hp drat wt vs am gear carb
## 5 2 mpg cyl disp hp drat wt vs am gear carb
## 5 3 mpg cyl disp hp drat wt vs am gear carb
## 5 4 mpg cyl disp hp drat wt vs am gear carb
## 5 5 mpg cyl disp hp drat wt vs am gear carb
```

```
head(mtcars_mice_rf)
```

```
##           mpg cyl disp  hp drat   wt  qsec vs am gear carb
## Mazda RX4      21.0   6  160 110 3.90 2.620 16.46  0  1    4    2
## Mazda RX4 Wag  21.0   6  160 110 3.90 2.875 17.02  0  1    5    1
## Datsun 710      22.8   4  108  93 3.85 2.320 18.61  1  1    4    1
## Hornet 4 Drive  21.4   6  258 110 3.08 3.215 19.44  1  0    3    1
## Hornet Sportabout 18.7   8  360 175 3.15 3.440 17.02  0  0    3    2
## Valiant        18.1   6  225 105 2.76 3.460 20.22  1  0    3    1
```

caret

caret is a full blown machine learning framework that contains a number of handy functions. `preprocess()` can be used for imputing, among other things. Here, we are using a `bagImpute`, which works by fitting a bagged tree model for each predictor (as a function of all the others). This method is quite accurate, but has much higher computational cost than say `medianImpute`.

```
library(caret)
set.seed(617)
mtcars_caret = predict(preProcess(mtcars_missing,
                                  method = 'bagImpute'),
                      newdata = mtcars_missing)

head(mtcars_caret)

##           mpg      cyl disp  hp drat   wt  qsec vs      am      gear
## Mazda RX4      21.0 6.000000  160 110 3.90 2.620 16.46  0 1.000000 4.000000
## Mazda RX4 Wag  21.0 6.000000  160 110 3.90 2.875 17.02  0 1.000000 3.615385
## Datsun 710      22.8 4.000000  108  93 3.85 2.320 18.61  1 1.000000 4.000000
## Hornet 4 Drive  21.4 6.000000  258 110 3.08 3.215 19.44  1 0.000000 3.000000
## Hornet Sportabout 18.7 6.246154  360 175 3.15 3.440 17.02  0 0.000000 3.615385
## Valiant        18.1 6.000000  225 105 2.76 3.460 20.22  1 0.3538462 3.000000
##           carb
## Mazda RX4      2.438462
## Mazda RX4 Wag  2.438462
## Datsun 710      1.000000
## Hornet 4 Drive  1.000000
## Hornet Sportabout 2.000000
## Valiant        2.438462
```

There are a number of imputation methods available but no single method is the best in every situation. Here is a comparison of the different imputation methods on the missing data.

```
df = data.frame(true_values = mtcars[is.na(mtcars_missing)],
               mice_pmm = mtcars_mice[is.na(mtcars_missing)],
               mice_rf = mtcars_mice_rf[is.na(mtcars_missing)],
               caret_bagImpute = mtcars_caret[is.na(mtcars_missing)])
```

```
df
```

```
##   true_values mice_pmm mice_rf caret_bagImpute
## 1      19.200    18.70  18.700      20.9530769
## 2      17.300    15.20  19.200      20.9530769
## 3      14.700    10.40  19.200      20.9530769
## 4      15.800    19.70  13.300      20.9530769
## 5      15.000    15.20  13.300      20.9530769
## 6       8.000     8.00   8.000       6.2461538
## 7       6.000     6.00   6.000       6.2461538
## 8       4.000     4.00   4.000       6.2461538
## 9     360.000   351.00 275.800     244.3669231
## 10    167.600   304.00 160.000     244.3669231
## 11    350.000   472.00 275.800     244.3669231
## 12    175.000   175.00 180.000     135.1461538
## 13    109.000    95.00 110.000     135.1461538
## 14      4.430     4.08   4.220       3.4791538
## 15      3.440     3.73   1.513       3.3680077
## 16      2.465     2.14   2.320       3.3680077
## 17      3.570     3.73   3.845       3.3680077
## 18      0.000     0.00   0.000       0.4307692
## 19      1.000     0.00   1.000       0.4307692
## 20      1.000     1.00   1.000       0.4307692
## 21      0.000     0.00   0.000       0.4307692
## 22      1.000     1.00   0.000       0.4307692
## 23      0.000     0.00   0.000       0.3538462
## 24      0.000     0.00   1.000       0.3538462
## 25      0.000     1.00   0.000       0.3538462
## 26      4.000     4.00   5.000       3.6153846
## 27      3.000     3.00   3.000       3.6153846
## 28      3.000     4.00   5.000       3.6153846
## 29      5.000     4.00   4.000       3.6153846
## 30      4.000     3.00   2.000       2.4384615
## 31      4.000     4.00   1.000       2.4384615
## 32      1.000     1.00   1.000       2.4384615
## 33      4.000     4.00   4.000       2.4384615
```

```
library(ggplot2); library(tidyr)
rownames(df)
```

```
## [1] "1" "2" "3" "4" "5" "6" "7" "8" "9" "10" "11" "12" "13" "14" "15"
## [16] "16" "17" "18" "19" "20" "21" "22" "23" "24" "25" "26" "27" "28" "29" "30"
## [31] "31" "32" "33"
```

```
df %>%
```

```
  mutate(id = as.integer(rownames(df)))%>%
```

```
  select(id, everything())%>%
```

```
  pivot_longer(2:5,names_to = 'method',values_to = 'values')%>%
```

```
  mutate(true_value = factor(ifelse(method == 'true_values',1,0),labels = c('True Value','Imputed Value')))
```

```
  ggplot(aes(x = id, y=values, color=true_value, shape = method))+
```

```
  geom_point()+theme_bw()
```

