

Advanced Trees with Wages Data

Contents

Data Description	1
Read Data	2
Split Data	2
Tree Models	2
Default Tree	2
Maximal Tree	2
Tree with Tuning	2
Ensemble Models	5
Bag Models	6
Random Forest Models	8
Tuned Random Forest	10
Forest with Ranger	11
Tuned Forest Ranger	11
Boosting Models	12
Boosting with cross-validation	13
Boosting with xgboost	13
Results	15
And the Winner is ???	16

Trees are flexible and easy to understand but tend to overfit and do not have the same level of predictive accuracy as other prediction models. In the following sections, we will examine a number of ways to improve tree models including tuning tree hyperparameters, and ensemble models such as bagging, forests, and boosting.

Data Description

Wages dataset is a simulated dataset based on a real dataset published in Data Analysis using Regression and Multilevel/Hierarchical Models by Andrew Gelman and Jennifer Hill

Data consists of characteristics of a set of employees and their annual earning. Variables included are:

- earn: Annual earning in dollars
- height: Height in inches
- gender: Gender (male, female)

- race: african-american, asian, hispanic, white
- ed: Years of education
- age: Age in years

Read Data

Set your working directory to the location where the data is saved by adapting the code below.

```
setwd('c:/my_classes/best_course_ever/') # on Windows
setwd('/my_classes/best_course_ever/') # on Mac

data = read.csv('wages.csv', stringsAsFactors = T)
data = data[data$earn>0,]
```

Split Data

```
set.seed(617)
split = sample(1:nrow(data),size = nrow(data)*0.8)
train = data[split,]
test = data[-split,]
```

In the following sections, we will build a series of models, in some cases tune them and then evaluate them on the test data.

Tree Models

In this section, we will examine a tree with default cp of 0.01, a maximal tree, and tree tuned with 10-fold cross-validation.

Default Tree

The default value of cp is 0.01

```
library(rpart); library(rpart.plot)
tree = rpart(earn~.,data=train)
pred = predict(tree,newdata=test)
rmse_tree = sqrt(mean((pred-test$earn)^2)); rmse_tree
## [1] 53456.15
```

Maximal Tree

This is the largest possible tree.

```
maximalTree = rpart(earn~.,data=train,control=rpart.control(cp=0))
pred = predict(maximalTree,newdata=test)
rmse_maximalTree = sqrt(mean((pred-test$earn)^2)); rmse_maximalTree
## [1] 57082.55
```

Tree with Tuning

Tune the complexity of a tree using 5-fold cross-validation. Here, we will examine cross-validation error for 100 different values of cp.

```

library(caret)
trControl = trainControl(method='cv',number = 5)
tuneGrid = expand.grid(.cp = seq(from = 0.001,to = 0.1,by = 0.001))
set.seed(617)
cvModel = train(earn~.,
                data=train,
                method="rpart",
                trControl = trControl,
                tuneGrid = tuneGrid)

cvModel$results

```

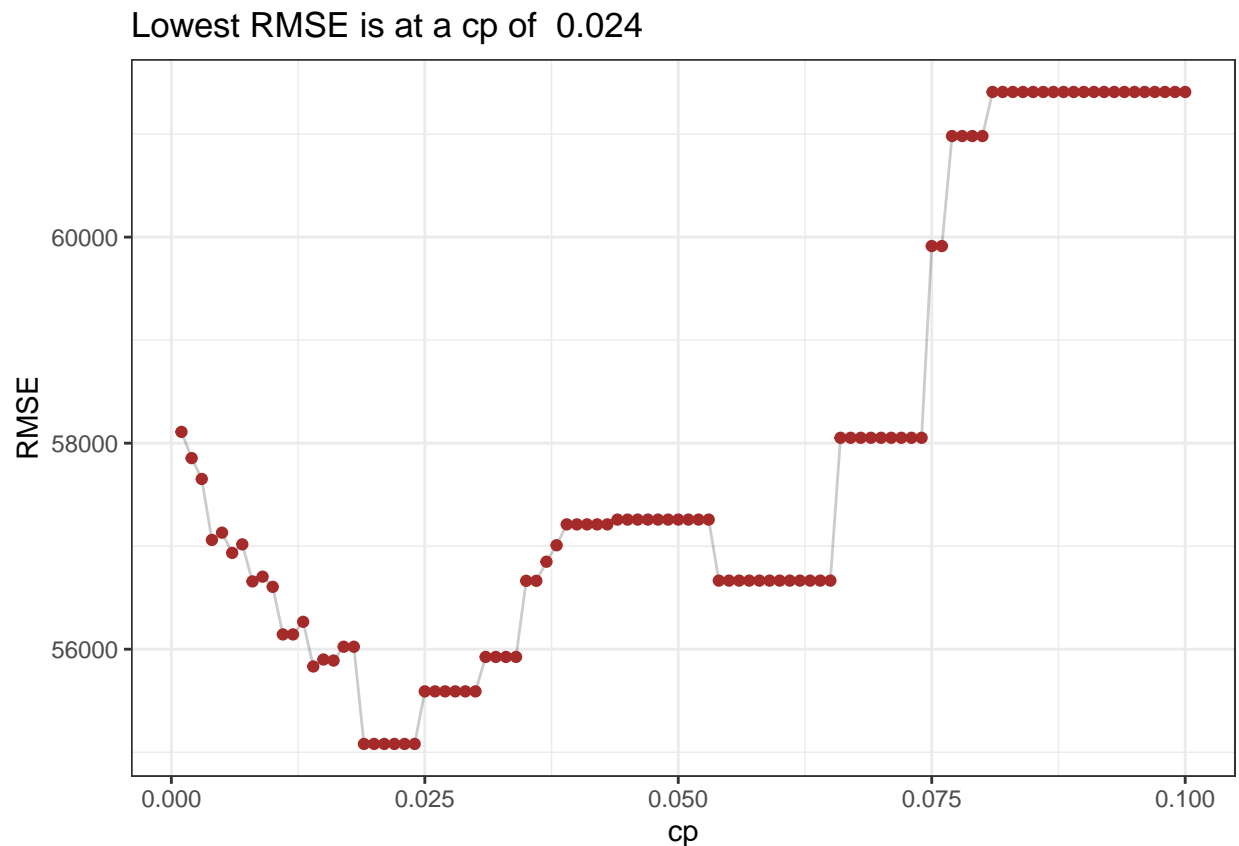
##	cp	RMSE	Rsquared	MAE	RMSESD	RsquaredSD	MAESD
## 1	0.001	58109.24	0.2205979	38289.82	4096.196	0.07289116	1409.1612
## 2	0.002	57854.68	0.2199231	38122.92	4255.585	0.08119947	1296.6183
## 3	0.003	57652.57	0.2201436	37787.52	4114.090	0.08668704	950.9004
## 4	0.004	57060.70	0.2261675	37474.47	4554.308	0.09144301	1385.1958
## 5	0.005	57131.09	0.2219554	37442.13	4730.652	0.08875365	1543.5813
## 6	0.006	56934.97	0.2236642	37421.03	4850.358	0.09632551	1327.4715
## 7	0.007	57017.74	0.2202879	37609.53	4789.538	0.09650538	1192.7217
## 8	0.008	56658.25	0.2274475	37378.90	4468.495	0.09804062	878.6218
## 9	0.009	56704.22	0.2259959	37435.16	4548.977	0.09891671	999.5121
## 10	0.010	56605.38	0.2257526	37498.57	4477.987	0.09881983	980.0040
## 11	0.011	56143.64	0.2352185	37513.93	3601.494	0.09873190	824.1015
## 12	0.012	56143.64	0.2352185	37513.93	3601.494	0.09873190	824.1015
## 13	0.013	56265.77	0.2323064	37818.42	3685.956	0.10019309	1178.7746
## 14	0.014	55832.55	0.2399879	37540.70	2844.735	0.09855232	627.1255
## 15	0.015	55900.22	0.2412936	37536.26	3021.938	0.10532164	623.4104
## 16	0.016	55889.92	0.2403748	37443.27	3023.800	0.10654123	476.9280
## 17	0.017	56023.99	0.2372423	37664.96	3336.998	0.11906822	1242.1867
## 18	0.018	56023.99	0.2372423	37664.96	3336.998	0.11906822	1242.1867
## 19	0.019	55080.56	0.2609847	37331.31	4357.008	0.14294082	1571.5555
## 20	0.020	55080.56	0.2609847	37331.31	4357.008	0.14294082	1571.5555
## 21	0.021	55080.56	0.2609847	37331.31	4357.008	0.14294082	1571.5555
## 22	0.022	55080.56	0.2609847	37331.31	4357.008	0.14294082	1571.5555
## 23	0.023	55080.56	0.2609847	37331.31	4357.008	0.14294082	1571.5555
## 24	0.024	55080.56	0.2609847	37331.31	4357.008	0.14294082	1571.5555
## 25	0.025	55590.83	0.2550573	37817.45	5127.188	0.15048022	2104.9636
## 26	0.026	55590.83	0.2550573	37817.45	5127.188	0.15048022	2104.9636
## 27	0.027	55590.83	0.2550573	37817.45	5127.188	0.15048022	2104.9636
## 28	0.028	55590.83	0.2550573	37817.45	5127.188	0.15048022	2104.9636
## 29	0.029	55590.83	0.2550573	37817.45	5127.188	0.15048022	2104.9636
## 30	0.030	55590.83	0.2550573	37817.45	5127.188	0.15048022	2104.9636
## 31	0.031	55925.05	0.2481645	38026.19	5699.440	0.16018328	2449.1204
## 32	0.032	55925.05	0.2481645	38026.19	5699.440	0.16018328	2449.1204
## 33	0.033	55925.05	0.2481645	38026.19	5699.440	0.16018328	2449.1204
## 34	0.034	55925.05	0.2481645	38026.19	5699.440	0.16018328	2449.1204
## 35	0.035	56663.93	0.2285600	38726.06	4903.507	0.15222101	2251.3395
## 36	0.036	56663.93	0.2285600	38726.06	4903.507	0.15222101	2251.3395
## 37	0.037	56849.17	0.2246281	38848.63	4937.458	0.14980035	2156.5893
## 38	0.038	57009.20	0.2223501	39155.68	4803.357	0.14711642	1713.7947
## 39	0.039	57211.47	0.2182940	39317.12	4873.200	0.14501027	1578.7061
## 40	0.040	57211.47	0.2182940	39317.12	4873.200	0.14501027	1578.7061
## 41	0.041	57211.47	0.2182940	39317.12	4873.200	0.14501027	1578.7061

## 42	0.042	57211.47	0.2182940	39317.12	4873.200	0.14501027	1578.7061
## 43	0.043	57211.47	0.2182940	39317.12	4873.200	0.14501027	1578.7061
## 44	0.044	57258.11	0.2208563	39407.07	4836.527	0.14816870	1460.3084
## 45	0.045	57258.11	0.2208563	39407.07	4836.527	0.14816870	1460.3084
## 46	0.046	57258.11	0.2208563	39407.07	4836.527	0.14816870	1460.3084
## 47	0.047	57258.11	0.2208563	39407.07	4836.527	0.14816870	1460.3084
## 48	0.048	57258.11	0.2208563	39407.07	4836.527	0.14816870	1460.3084
## 49	0.049	57258.11	0.2208563	39407.07	4836.527	0.14816870	1460.3084
## 50	0.050	57258.11	0.2208563	39407.07	4836.527	0.14816870	1460.3084
## 51	0.051	57258.11	0.2208563	39407.07	4836.527	0.14816870	1460.3084
## 52	0.052	57258.11	0.2208563	39407.07	4836.527	0.14816870	1460.3084
## 53	0.053	57258.11	0.2208563	39407.07	4836.527	0.14816870	1460.3084
## 54	0.054	56665.90	0.2307355	39257.96	4936.649	0.13656457	1366.5363
## 55	0.055	56665.90	0.2307355	39257.96	4936.649	0.13656457	1366.5363
## 56	0.056	56665.90	0.2307355	39257.96	4936.649	0.13656457	1366.5363
## 57	0.057	56665.90	0.2307355	39257.96	4936.649	0.13656457	1366.5363
## 58	0.058	56665.90	0.2307355	39257.96	4936.649	0.13656457	1366.5363
## 59	0.059	56665.90	0.2307355	39257.96	4936.649	0.13656457	1366.5363
## 60	0.060	56665.90	0.2307355	39257.96	4936.649	0.13656457	1366.5363
## 61	0.061	56665.90	0.2307355	39257.96	4936.649	0.13656457	1366.5363
## 62	0.062	56665.90	0.2307355	39257.96	4936.649	0.13656457	1366.5363
## 63	0.063	56665.90	0.2307355	39257.96	4936.649	0.13656457	1366.5363
## 64	0.064	56665.90	0.2307355	39257.96	4936.649	0.13656457	1366.5363
## 65	0.065	56665.90	0.2307355	39257.96	4936.649	0.13656457	1366.5363
## 66	0.066	58052.07	0.1887969	39568.77	5076.402	0.11307757	1167.4040
## 67	0.067	58052.07	0.1887969	39568.77	5076.402	0.11307757	1167.4040
## 68	0.068	58052.07	0.1887969	39568.77	5076.402	0.11307757	1167.4040
## 69	0.069	58052.07	0.1887969	39568.77	5076.402	0.11307757	1167.4040
## 70	0.070	58052.07	0.1887969	39568.77	5076.402	0.11307757	1167.4040
## 71	0.071	58052.07	0.1887969	39568.77	5076.402	0.11307757	1167.4040
## 72	0.072	58052.07	0.1887969	39568.77	5076.402	0.11307757	1167.4040
## 73	0.073	58052.07	0.1887969	39568.77	5076.402	0.11307757	1167.4040
## 74	0.074	58052.07	0.1887969	39568.77	5076.402	0.11307757	1167.4040
## 75	0.075	59913.26	0.1468932	40537.95	4585.438	0.09474611	2372.1906
## 76	0.076	59913.26	0.1468932	40537.95	4585.438	0.09474611	2372.1906
## 77	0.077	60981.05	0.1152880	41022.86	4915.524	0.06462534	2092.4692
## 78	0.078	60981.05	0.1152880	41022.86	4915.524	0.06462534	2092.4692
## 79	0.079	60981.05	0.1152880	41022.86	4915.524	0.06462534	2092.4692
## 80	0.080	60981.05	0.1152880	41022.86	4915.524	0.06462534	2092.4692
## 81	0.081	61408.86	0.1011505	41432.09	4299.847	0.07159193	1957.6928
## 82	0.082	61408.86	0.1011505	41432.09	4299.847	0.07159193	1957.6928
## 83	0.083	61408.86	0.1011505	41432.09	4299.847	0.07159193	1957.6928
## 84	0.084	61408.86	0.1011505	41432.09	4299.847	0.07159193	1957.6928
## 85	0.085	61408.86	0.1011505	41432.09	4299.847	0.07159193	1957.6928
## 86	0.086	61408.86	0.1011505	41432.09	4299.847	0.07159193	1957.6928
## 87	0.087	61408.86	0.1011505	41432.09	4299.847	0.07159193	1957.6928
## 88	0.088	61408.86	0.1011505	41432.09	4299.847	0.07159193	1957.6928
## 89	0.089	61408.86	0.1011505	41432.09	4299.847	0.07159193	1957.6928
## 90	0.090	61408.86	0.1011505	41432.09	4299.847	0.07159193	1957.6928
## 91	0.091	61408.86	0.1011505	41432.09	4299.847	0.07159193	1957.6928
## 92	0.092	61408.86	0.1011505	41432.09	4299.847	0.07159193	1957.6928
## 93	0.093	61408.86	0.1011505	41432.09	4299.847	0.07159193	1957.6928
## 94	0.094	61408.86	0.1011505	41432.09	4299.847	0.07159193	1957.6928
## 95	0.095	61408.86	0.1011505	41432.09	4299.847	0.07159193	1957.6928

```
## 96 0.096 61408.86 0.1011505 41432.09 4299.847 0.07159193 1957.6928
## 97 0.097 61408.86 0.1011505 41432.09 4299.847 0.07159193 1957.6928
## 98 0.098 61408.86 0.1011505 41432.09 4299.847 0.07159193 1957.6928
## 99 0.099 61408.86 0.1011505 41432.09 4299.847 0.07159193 1957.6928
## 100 0.100 61408.86 0.1011505 41432.09 4299.847 0.07159193 1957.6928
```

Results from Tuning

```
library(ggplot2)
ggplot(data=cvModel$results, aes(x=cp, y=RMSE))+
  geom_line(size=0.5,alpha=0.2)+
  geom_point(color='brown')+
  theme_bw()+
  ggtitle(label=paste('Lowest RMSE is at a cp of ',cvModel$bestTune$cp))
```



Evaluate Tuned model on Test sample

```
cvTree = rpart(earn~.,data=train,cp = cvModel$bestTune$cp)
pred = predict(cvTree,newdata=test)
rmse_cvTree = sqrt(mean((pred-test$earn)^2)); rmse_cvTree
## [1] 54198.35
```

Ensemble Models

- Bagging
- Forest
- Boosting

Bag Models

Bootstrap Aggregation models generate a large number of bootstrapped samples. A tree is fit to each bootstrapped sample. Predictions are generating as an average of all models (for numerical outcome variables) or the majority group (for categorical outcome variables).


These models can be implemented using many packages including `ipred`, `randomForest`, `adabag`, `bagEarth`, `treeBag`, `bagFDA`. In this illustration, we are using a `randomForest` model by setting `mtry` to be the number of predictors, i.e., 5.

```
library(randomForest)
set.seed(617)
bag = randomForest(earn~.,data=train,mtry = ncol(train)-1,ntree=1000)
pred = predict(bag,newdata=test)
rmse_bag = sqrt(mean((pred-test$earn)^2)); rmse_bag

## [1] 53486.83
```

The various trees that were fit vary in their size


```
hist(treesize(bag))
```



```
# getTree(bag,k=100) # One can examine a given tree
```

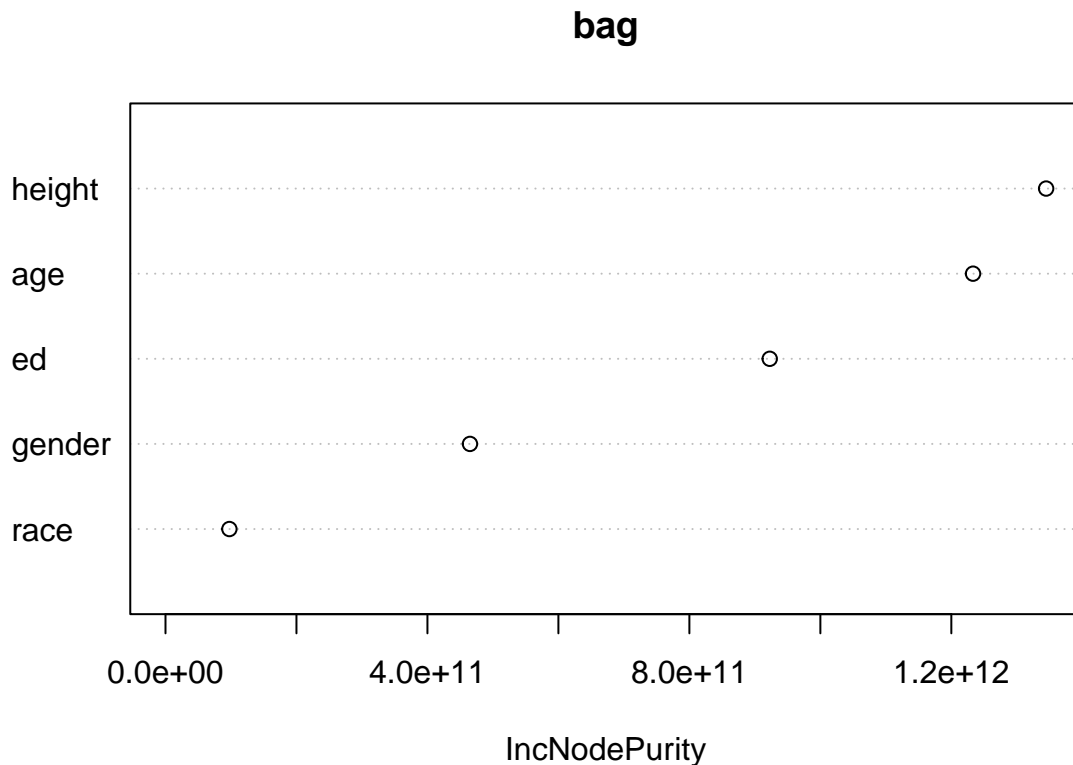
Generally speaking, error for a bag model decreases with an increase in the trees.

```
plot(bag)
```



Relative importance of predictors

```
varImpPlot(bag)
```



```
importance(bag)
```

```
##          IncNodePurity
## height  1.344749e+12
## gender  4.649394e+11
## race    9.759492e+10
## ed      9.227044e+11
## age     1.233138e+12
```


Random Forest Models

Unlike bag models which consider all predictors in constructing each tree, randomForest models consider a subset of predictors (default is $p/3$ for numerical outcomes or \sqrt{p} for categorical outcomes) for constructing each tree. In `library(randomForest)`, `mtry` controls number of variables considered for each tree.

```
library(randomForest)
set.seed(617)
forest = randomForest(earn~.,data=train,ntree = 1000)
pred = predict(forest,newdata=test)
rmse_forest = sqrt(mean((pred-test$earn)^2)); rmse_forest
## [1] 50062.28
```


As in the case for bags, more trees reduces error to a certain point, after which the graph asymptotes.

```
plot(forest)
```

Relative importance of predictors

```
varImpPlot(forest)
```



Tuned Random Forest

The `mtry` parameter of a `randomForest` model can be tuned to improve model predictions. Here, we will use 5-fold cross validation to examine four values of `mtry`.

```
trControl=trainControl(method="cv",number=5)
tuneGrid = expand.grid(mtry=1:4)
set.seed(617)
cvModel = train(earn~.,data=train,
                 method="rf",ntree=1000,trControl=trControl,tuneGrid=tuneGrid )
cvModel
```

```
## Random Forest
##
## 1094 samples
##    5 predictor
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 875, 875, 875, 876, 875
## Resampling results across tuning parameters:
##
##  mtry  RMSE      Rsquared  MAE
##  1     55791.64  0.2891217  36982.68
##  2     53265.49  0.3015901  35135.91
##  3     54021.84  0.2856715  35466.80
##  4     55117.06  0.2680037  36034.66
```

```
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was mtry = 2.

Now, let us use the best mtry of 2.

cvForest = randomForest(earn~.,data=train,ntree = 1000,mtry=cvModel$bestTune$mtry)
pred = predict(cvForest,newdata=test)
rmse_cv_forest = sqrt(mean((pred-test$earn)^2)); rmse_cv_forest
## [1] 52363.92
```

Forest with Ranger

ranger is popular R library for running random forest models.

```
library(ranger)

##
## Attaching package: 'ranger'

## The following object is masked from 'package:randomForest':
##
##      importance

forest_ranger = ranger(earn~.,data=train,num.trees = 1000)
pred = predict(forest_ranger, data =test,num.trees = 1000)
rmse_forest_ranger = sqrt(mean((pred$predictions-test$earn)^2)); rmse_forest_ranger
## [1] 52381.18
```

Tuned Forest Ranger

Tuning ranger model for different values of mtry, splitrule and min.node.size

```
trControl=trainControl(method="cv",number=5)
tuneGrid = expand.grid(mtry=1:4,
                      splitrule = c('variance','extratrees','maxstat'),
                      min.node.size = c(2,5,10,15,20,25))

set.seed(617)
cvModel = train(earn~.,
               data=train,
               method="ranger",
               num.trees=1000,
               trControl=trControl,
               tuneGrid=tuneGrid )
cv_forest_ranger = ranger(earn~.,
                        data=train,
                        num.trees = 1000,
                        mtry=cvModel$bestTune$mtry,
                        min.node.size = cvModel$bestTune$min.node.size,
                        splitrule = cvModel$bestTune$splitrule)
pred = predict(cv_forest_ranger, data =test, num.trees = 1000)
rmse_cv_forest_ranger = sqrt(mean((pred$predictions-test$earn)^2)); rmse_cv_forest_ranger
## [1] 51741.44
```


Boosting Models

Like bag and forest models, boosting models are ensemble models that derive predictions from a number of trees. The key difference is that in boosting, trees are grown sequentially, each tree is grown using information from previously grown trees. Thus, boosting can be seen as a slow learning evolutionary model. Since we are predicting a numerical variable, `earn`, the distribution is set to 'gaussian'. Had the goal been to predict a binary outcome, we would have set distribution to 'bernoulli'.

```
library(gbm)
set.seed(617)
boost = gbm(earn~.,
            data=train,
            distribution="gaussian",
            n.trees = 500,
            interaction.depth = 2,
            shrinkage = 0.01)
```

Boosting models offer a way to examine the relative influence of variables.

```
summary(boost)
```



```
##      var      rel.inf
## ed      ed 41.85723381
## age     age 23.85348758
## gender  gender 19.82020900
## height height 14.41961849
## race    race  0.04945111
```

Let us examine RMSE for the train sample.

```

pred = predict(boost,n.trees = 500)
rmse_boost_train = sqrt(mean((pred-train$earn)^2)); rmse_boost_train
## [1] 50425.2

Now, examine RMSE for test sample.

pred = predict(boost,newdata=test,n.trees = 500)
rmse_boost = sqrt(mean((pred-test$earn)^2)); rmse_boost
## [1] 50876.59

```

Boosting with cross-validation

Boosting models are notorious for overfitting training data. A very simple way to see this borne out is to see the effect of increasing number of trees on train and test rmse. Specifically, in the model above, try running the models with the following number of trees: 200, 1e3, 1e4, 1e6. To avoid the folly of overfitting, it is best to tune the model using cross-validation. In the code that follows, we will tune a gradient boosting model using interaction depth, shrinkage and minobsinnode.

The code below uses a garbage collector to hide from us the long list of models tested.

```

library(caret)
set.seed(617)
trControl = trainControl(method="cv",number=5)
tuneGrid = expand.grid(n.trees = 500,
                      interaction.depth = c(1,2,3),
                      shrinkage = (1:100)*0.001,
                      n.minobsinnode=c(5,10,15))
garbage = capture.output(cvModel <- train(earn~.,
                                         data=train,
                                         method="gbm",
                                         trControl=trControl,
                                         tuneGrid=tuneGrid))

set.seed(617)
cvBoost = gbm(earn~.,
              data=train,
              distribution="gaussian",
              n.trees=cvModel$bestTune$n.trees,
              interaction.depth=cvModel$bestTune$interaction.depth,
              shrinkage=cvModel$bestTune$shrinkage,
              n.minobsinnode = cvModel$bestTune$n.minobsinnode)
pred = predict(cvBoost,test,n.trees=500)
rmse_cv_boost = sqrt(mean((pred-test$earn)^2)); rmse_cv_boost
## [1] 51234.84

```

Boosting with xgboost

XGBoost is an optimized distributed gradient boosting library designed to be highly efficient, flexible and portable.

The algorithm is also a bit picky about the format of variables used. All factor class variables need to be dummy coded and fed into the model as a matrix. To do this, we will dummy code using library(vtreat)

```

library(vtreat)

## Loading required package: wrapr

```

```

trt = designTreatmentsZ(dframe = train,
                        varlist = names(train)[2:6])

## [1] "vtreat 1.6.3 inspecting inputs Wed Sep 01 17:09:53 2021"
## [1] "designing treatments Wed Sep 01 17:09:53 2021"
## [1] " have initial level statistics Wed Sep 01 17:09:53 2021"
## [1] " scoring treatments Wed Sep 01 17:09:53 2021"
## [1] "have treatment plan Wed Sep 01 17:09:53 2021"

newvars = trt$scoreFrame[trt$scoreFrame$code%in% c('clean','lev'),'varName']

train_input = prepare(treatmentplan = trt,
                      dframe = train,
                      varRestriction = newvars)
test_input = prepare(treatmentplan = trt,
                     dframe = test,
                     varRestriction = newvars)

head(train_input)

##   height ed age gender_lev_x_female gender_lev_x_male
## 1  65.98 14  47                1                0
## 2  64.07 16  64                1                0
## 3  59.61 16  92                1                0
## 4  63.28 12  42                1                0
## 5  71.34 11  22                0                1
## 6  67.70 14  29                1                0
##   race_lev_x_african_minus_american race_lev_x_asian race_lev_x_hispanic
## 1                0                0                0
## 2                0                0                0
## 3                0                1                0
## 4                0                0                0
## 5                1                0                0
## 6                0                0                0
##   race_lev_x_white
## 1                1
## 2                1
## 3                0
## 4                1
## 5                0
## 6                1

```


Like all boosting models, xgboost can overfit the train data. To identify the optimal rounds, we use a cross-validation function within xgboost.

```

library(xgboost); library(caret)
set.seed(617)
tune_nrounds = xgb.cv(data=as.matrix(train_input),
                      label = train$earn,
                      nrounds=250,
                      nfold = 5,
                      verbose = 0)

ggplot(data=tune_nrounds$evaluation_log, aes(x=iter, y=test_rmse_mean))+
  geom_point(size=0.4, color='sienna')+
  geom_line(size=0.1, alpha=0.1)+
  theme_bw()

```



As is obvious from the above chart, the optimal nrounds is a rather small number

```
which.min(tune_nrounds$evaluation_log$test_rmse_mean)
```

```
## [1] 6
```

Next, we use xgboost to fit the train data with nrounds = 6 and apply the model to the test data.

```
xgboost2= xgboost(data=as.matrix(train_input),
                  label = train$earn,
                  nrounds=6,
                  verbose = 0)
pred = predict(xgboost2,
               newdata=as.matrix(test_input))
rmse_xgboost = sqrt(mean((pred - test$earn)^2)); rmse_xgboost
## [1] 51597.19
```

Now, it is possible to tune an xgboost model using other parameters such as eta, max_depth, and gamma using the train function from library(caret)

Results

Here are the results from the different models that were run

```
data.frame(models = c('Tree', 'Maximal Tree', 'Tuned Tree', 'Bag', 'Forest', 'Tuned Forest', 'Ranger', 'Tuned
                    RMSE = c(rmse_tree, rmse_maximalTree, rmse_cvTree, rmse_bag, rmse_forest, rmse_cv_forest, rmse_cv_forest_ranger, rmse_boost, rmse_cv_boost, rmse_xgboost))
##          models      RMSE
```

```
## 1      Tree 53456.15
## 2 Maximal Tree 57082.55
## 3   Tuned Tree 54198.35
## 4      Bag 53486.83
## 5   Forest 50062.28
## 6 Tuned Forest 52363.92
## 7   Ranger 52381.18
## 8 Tuned Ranger 51741.44
## 9   Boost 50876.59
## 10 Tuned Boost 51234.84
## 11   XGBoost 51597.19
```

And the Winner is ???